

# K Nearest Neighbors (KNN) Cheat Sheet

Mastering Classification & Regression: Concepts, Math, and Implementation

## 1 What is K Nearest Neighbors

**Definition:** A non-parametric, supervised learning algorithm that classifies or predicts a value for a new data point based on the  $K$  closest training data points.

**How it Works:**


1. Store all training data.
2. For a new input, calculate distance to all training points.
3. Select the  $K$  nearest points.
4. **Vote** (Classification) or **Average** (Regression).

**Lazy Learning:** KNN does not "learn" a model (weights/biases) during training. It simply stores data. All computation happens at prediction time (Eager evaluation).

**Use Cases:** Recommendation systems, pattern recognition, data imputation.


## 2 KNN Classification vs Regression

**Classification (Voting):**



- **Output:** Discrete Class Label.
- **Mechanism:** Majority Vote (Mode).
- **Example:** Is this email Spam or Not?


**Regression (Averaging):**



- **Output:** Continuous Value.
- **Mechanism:** Average (Mean) of neighbors' values.
- **Example:** Estimate house price based on 3 neighbors.


## 3 Distance Metrics

How "closeness" is measured drastically affects performance.



- **Euclidean (L2):** Straight line. Standard for continuous data. Sensitive to scale.
- **Manhattan (L1):** Grid path. Good for high dimensions or sparse data.
- **Minkowski:** Generalization.  $p = 1$  is Manhattan,  $p = 2$  is Euclidean.
- **Cosine:** Angle between vectors. Best for text/NLP where magnitude matters less than direction.


## 4 Choosing the Value of K



- **Small K (e.g., 1):** Low Bias, High Variance. Model captures noise. \*\*Overfitting\*\*. Decision boundary is jagged.
- **Large K:** High Bias, Low Variance. Model smooths over details. \*\*Underfitting\*\*. Computationally cheaper.
- **Selection:** Use \*\*Cross-Validation\*\* (e.g., Elbow Method) to find the  $K$  that minimizes validation error.

## 5 Data Preprocessing

**Feature Scaling (CRITICAL):** KNN calculates distances. If Feature A ranges [0-1] and Feature B [0-1000], Feature B will dominate the distance calculation.



**Standardization:**  $\frac{x-\mu}{\sigma}$  (Mean 0, Var 1). Good for outliers.

**Normalization:**  $\frac{x-min}{max-min}$  (Range 0-1). Good for fixed bounds.

**Handling Categorical:** Must convert to numeric (One-Hot Encoding) as distance functions require numbers.

## 6 KNN Algorithm Step-by-Step

1. **Load Data:** Initialize training set  $D$ .
2. **Input:** Receive query point  $q$ .
3. **Distance:** For every point  $x_i$  in  $D$ , calculate  $dist(q, x_i)$ .
4. **Sort:** Order points by increasing distance.
5. **Select:** Pick top  $K$  points.
6. **Aggregate:** - **Class:** Return Mode(Labels). - **Reg:** Return Mean(Values).

## 7 End-to-End Classification

**Dataset:** Iris (Predict species).

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

# 1. Load & Split
data = load_iris()
X_train, X_test, y_train, y_test = train_test_split(
    data.data, data.target, test_size=0.2
)

# 2. Scale (Mandatory)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# 3. Train
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)

# 4. Predict & Evaluate
preds = knn.predict(X_test)
print(f"Acc: {accuracy_score(y_test, preds)}")
```

## 8 End-to-End Regression

**Dataset:** California Housing (Predict price).

```
from sklearn.datasets import fetch_california_housing
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error

# 1. Load & Split
data = fetch_california_housing()
X_train, X_test, y_train, y_test = train_test_split(
    data.data, data.target, test_size=0.2
)

# 2. Scale
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# 3. Train
knn_reg = KNeighborsRegressor(n_neighbors=5)
knn_reg.fit(X_train, y_train)

# 4. Predict & Evaluate
preds = knn_reg.predict(X_test)
mse = mean_squared_error(y_test, preds)
print(f"RMSE: {mse**0.5:.2f}")
```

## 9 Performance Metrics

**Classification:**

- **Accuracy:** Correct / Total.
- **Precision/Recall:** Quality of positive predictions.
- **Confusion Matrix:** Shows Type I/II errors.

**Regression:**

- **MSE:** Mean Squared Error (Penalizes large errors).
- **RMSE:** Root MSE (Same units as target).
- **MAE:** Mean Absolute Error (Robust to outliers).

## 10 Hyperparameter Tuning

Using GridSearchCV to find the best model.


- **n\_neighbors:** Range [1, 30]. Odd numbers preferred to avoid ties.
- **weights:** 'uniform' (all neighbors equal) vs 'distance' (closer neighbors count more).
- **metric:** 'euclidean', 'manhattan'.

## 11 Complexity & Scalability

**Time Complexity:**

- Training:  $O(1)$  (Just storage).
- Prediction:  $O(N \times D)$  where  $N$  is samples,  $D$  is dimensions. **Very Slow** for large  $N$ .

**Scalability Issues:** KNN requires calculating distance to \*every\* point. **Solution:** Use Approximate Nearest Neighbors (ANN) or Tree structures:



## 12 Pitfalls & Best Practices

**Curse of Dimensionality:** As features ( $D$ ) increase, data becomes sparse. Distances between "nearest" and "farthest" neighbors converge, making distance meaningless. \*Fix\*: Use Dimensionality Reduction (PCA) before KNN.

**Irrelevant Features:** Noise features dilutes distance. \*Fix\*: Feature Selection.

**Imbalanced Data:** Majority class dominates voting. \*Fix\*: Use 'weights='distance'".

## 13 Summary

Concept	Key Takeaway
Type	Instance-based, Lazy
Metric	Euclidean (Standard), Cosine (Text)
Scaling	<b>Mandatory</b>
K	Hyperparameter (Tune it!)
Speed	Fast Train, Slow Predict