

# Naive Bayes: Regression & Classification Cheat Sheet

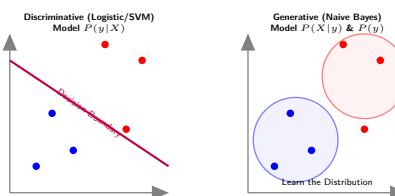
Probabilistic Machine Learning: Theory, Implementation, and Best Practices

## 1 Introduction

**Definition:** A family of probabilistic algorithms based on applying Bayes' theorem with a strong ("naive") assumption of independence between features.

**Why "Naive"?** It assumes all features are independent of each other given the class label. In reality, features are often correlated (e.g., "San" and "Francisco"), but NB ignores this to simplify computation.

**Model Type:**



- Generative:** Models how the data is generated ( $P(X|y)$ ). Uses prior probability.
- Discriminative:** Models the boundary directly ( $P(y|X)$ ).

## 2 Mathematical Foundation

$$P(y|X) = \frac{\overbrace{P(X|y) \cdot P(y)}^{\text{Posterior}}}{\underbrace{P(X)}_{\text{Evidence}}}$$

**Core Formula:**

$$P(y|x_1, \dots, x_n) = \frac{P(y) \prod_{i=1}^n P(x_i|y)}{P(X)}$$

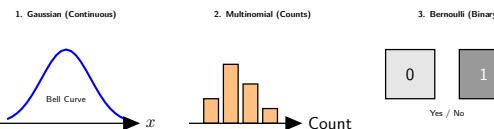
**Components:**

- Prior  $P(y)$ :** Probability of class  $y$  before seeing data.
- Likelihood  $P(x_i|y)$ :** Probability of feature  $x_i$  appearing in class  $y$ .
- Posterior  $P(y|X)$ :** The probability we want to calculate.

**Decision Rule:**

$$\hat{y} = \arg \max_y P(y) \prod_{i=1}^n P(x_i|y)$$

## 3 Types of Naive Bayes



### 1. Gaussian NB:

- Use Case:** Continuous features (e.g., Iris data, sensor readings).

- Assumption:** Features follow a Normal (Gaussian) distribution.

### 2. Multinomial NB:

- Use Case:** Discrete counts (e.g., Word counts in text).
- Assumption:** Data follows a Multinomial distribution.

### 3. Bernoulli NB:

- Use Case:** Binary/Boolean features (e.g., Word presence/absence).

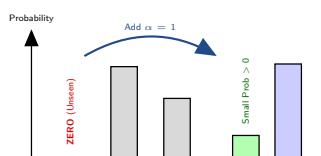
- Assumption:** Data is binary (0/1).

## 4 Data Requirements

**Features:**

- Continuous:** Use GaussianNB.
- Categorical:** Convert to counts/One-Hot and use MultinomialNB.
- Text:** Bag of Words or TF-IDF.

**Laplace Smoothing (Additive Smoothing):** Solves the "Zero Probability" problem. If a word never appears in training, Likelihood = 0, killing the whole probability product.



$$P(x_i|y) = \frac{\text{count}(x_i, y) + \alpha}{\text{count}(y) + \alpha \cdot N_{\text{features}}}$$

where  $\alpha = 1$  is standard Laplace smoothing.

## 5 Classification Implementation

**Scenario:** Spam Detection (Text).

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline

# 1. Data
X = ["offer is secret", "click secret link", "meeting today"]
y = [1, 1, 0] # 1=Spam, 0=Ham

# 2. Split
X_train, X_test, y_train, y_test = train_test_split(X, y)

# 3. Pipeline (Vectorize -> Model)
model = make_pipeline(
    CountVectorizer(), # Convert text to counts
    MultinomialNB(alpha=1.0) # Naive Bayes
)

# 4. Train & Predict
model.fit(X_train, y_train)
preds = model.predict(X_test)
print(preds)
```

## 6 Regression with Bayes

Standard Naive Bayes is a **Classifier**. For Regression, we use **Bayesian Ridge** or **Gaussian Process**. It assumes the target  $y$  comes from a probability distribution.

**Bayesian Ridge:** Assumes  $y$  is Gaussian distributed around  $Xw$ . It estimates a distribution over the weights  $w$ , not just point estimates.

```
from sklearn.linear_model import BayesianRidge

# 1. Data
X = [[0, 0], [1, 1], [2, 2]]
y = [0, 1, 2]

# 2. Model
reg = BayesianRidge()

# 3. Train
reg.fit(X, y)

# 4. Predict (Returns Mean and Std Dev)
mean, std = reg.predict([[1, 0]], return_std=True)
print(f"Pred: {mean:.2f} +/- {std:.2f}")
```

## 7 Performance Metrics

**Classification:**

Actual		
Predicted	TP (Hit)	FP (Type I)
FN (Type II)		TN (Rejection)

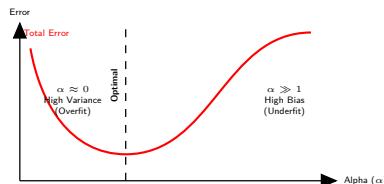
- **Accuracy:** % Correct.
- **Precision:**  $P(\text{True Spam} \mid \text{Pred Spam})$ . Critical for spam filters.
- **Recall:**  $P(\text{Pred Spam} \mid \text{True Spam})$ .
- **Log Loss:** Penalty for confident wrong answers.

#### Regression:

- **MSE:** Mean Squared Error.
- **R2 Score:** Variance explained.

## 8 Tuning

Smoothing ( $\alpha$ ):



- $\alpha = 0$ : No smoothing (Risk of zero prob).
- $\alpha = 1$ : Laplace smoothing (Default).
- $\alpha > 1$ : High smoothing (High Bias, Low Variance).

**Priors:** Can explicitly set 'priors=[0.2, 0.8]' if you know the class balance differs from the training set.

## 9 Practical Use Cases

- **Spam Filtering:** The classic use case. Fast, handles high-dim text well.
- **Sentiment Analysis:** Positive/Negative classification.
- **Real-time Prediction:** Extremely fast inference speed.
- **Baseline Model:** Always run NB first to set a baseline for complex models.

## 10 Common Mistakes

- **Correlated Features:** "New" and "York" in "New York" are correlated. NB counts them twice, inflating confidence.
- **Zero Probability:** Forgetting smoothing ( $\alpha$ ) causes the model to crash on unseen words.
- **Wrong Distribution:** Using GaussianNB on sparse word counts (Use Multinomial instead).

## 11 Comparison

### Model vs Naive Bayes

Logistic Reg	Discriminative. Slower. Handles correlation better.
SVM	Discriminative. Max margin. Slower. Better acc.
Trees	Non-linear. Handles interactions. Much slower.

## 12 Summary

- **Choose When:** High dimensions (Text), small data, need speed, need baseline.
- **Avoid When:** Features are highly correlated, huge training data available (DL might be better).
- **Key Param:** Alpha ( $\alpha$ ) for smoothing.