

C++ Allocators

Why Allocators Exist

Imagine you're using:

```
std::vector<int> v;  
v.push_back(1);  
v.push_back(2);
```

When you create a container like this, it uses the default allocator, which internally calls the global heap allocator (`::operator new()` / `malloc()`).

That means every time the vector grows, it may request new memory from the OS heap — which can be slow, fragmented, or unpredictable. That's fine for normal programs...

But in low-latency or high-performance systems (e.g. trading, game engines, or networking frameworks), this is a big problem:

- Every `malloc/new` call can trigger mutex locks, heap fragmentation, and cache misses.
- Each allocation can cost hundreds of nanoseconds to microseconds — unacceptable when your SLA is <10 µs.

 Allocators solve this by letting you control where and how memory is managed.

C++ Allocators

The Core Idea

→ An allocator is a pluggable memory management strategy for STL containers.

Allocators let you change that memory source.

For example:

```
std::pmr::monotonic_buffer_resource pool;
std::pmr::vector<int> v{ &pool };
```

Here's what happens:

- You first create a pool (a pre-allocated block of memory).
- Then you tell the vector to use that pool instead of the heap.

Now:

- ✓ Every allocation comes from the pool, not the global heap.
- ✓ Allocation is just a pointer bump inside that pre-allocated memory.
- ✓ No locks, no system calls, no fragmentation.
- ✓ When done, you can free the whole pool in one go — no per-object frees.

```
Container → Allocator → Memory Pool
(instead of Heap)
```

C++ Allocators

Allocators separate **data structure logic** from **memory management logic**.

This is why STL containers can use **any allocator** you give them:

```
| std::vector<int, CustomAllocator<int>> vec;
```

or

```
| std::pmr::vector<int> vec{&resource};
```

C++ provides two key building blocks for memory management:

`std::allocator`

`std::allocator_traits`

`std::allocator`

The default allocator in the C++ Standard Library; it obtains and releases memory from the global heap using operator new and operator delete.

`std::allocator_traits`

A wrapper that provides a uniform interface for all allocator types, letting STL containers interact generically with different allocator implementations.

C++ Allocators

Program 1 — Using `std::allocator` directly

```
#include <iostream>
#include <memory>

int main() {
    std::allocator<int> alloc;

    // allocate space for 3 ints
    int* p = alloc.allocate(3);

    for (int i = 0; i < 3; ++i) {
        alloc.construct(&p[i], i * 10);
    }

    for (int i = 0; i < 3; ++i) {
        std::cout << p[i] << " ";
    }

    std::cout << "\n";

    for (int i = 0; i < 3; ++i) {
        alloc.destroy(&p[i]);
    }

    alloc.deallocate(p, 3);
}
```

C++ Allocators

Program 2 — Using `std::allocator_traits` directly

```
#include <iostream>
#include <memory>

int main() {
    std::allocator<int> alloc;
    using Traits =
        std::allocator_traits<decltype(alloc)>;

    int* p = Traits::allocate(alloc, 3);
    for (int i = 0; i < 3; ++i) {
        Traits::construct(alloc, &p[i], i + 1);
    }

    for (int i = 0; i < 3; ++i) {
        std::cout << p[i] << " ";
    }

    std::cout << "\n";

    for (int i = 0; i < 3; ++i) {
        Traits::destroy(alloc, &p[i]);
    }

    Traits::deallocate(alloc, p, 3);
}
```