

Linux Mastery for C++ Developers

A Comprehensive Roadmap with Commands, Tools, and Real-World Examples

by CompilerSutra

Version 3.1 | November 2, 2025



“The Linux philosophy is ‘Laugh in the face of danger’. Oops. Wrong One. ‘Do it yourself’. Yes, that’s it.”

— Linus Torvalds

Contents

1	Introduction	2
1.1	Why Linux Matters for C++ Developers	2
1.2	How to Use This Roadmap	2
1.3	Real-world Use Cases	2
2	Linux Basics (Beginner Level)	3
2.1	Filesystem Hierarchy Standard	3
2.2	File Operations Mastery	3
2.3	Advanced Text Editing	3
2.4	Permission Deep Dive	4
3	Intermediate Linux Commands	5
3.1	Advanced File Searching	5
3.2	System Performance Monitoring	5
3.3	Network Diagnostics	5
3.4	Package Management Deep Dive	6
4	Advanced System Administration	7
4.1	Systemd Service Management	7
4.2	Kernel and System Optimization	7
4.3	Advanced Security Features	7
4.4	Containerization with Podman/Docker	8
5	Shell Scripting Excellence	9
5.1	Advanced Bash Features	9
5.2	Error Handling and Debugging	9
5.3	Configuration Management	9
6	Text Processing and Data Manipulation	11
6.1	grep, rg, awk, sed	11
6.2	cut, sort, uniq, tr, wc	11
6.3	Practical assembly / log parsing	11
7	Developer Productivity Tools	12
7.1	tmux Essentials	12
7.2	fzf, ripgrep, bat	12
7.3	Editor setups: clangd, LSP, ccls	12
8	Build Systems & Source Control	13
8.1	CMake advanced patterns	13
8.2	Toolchain files and cross-compiling	13
8.3	Makefile tips and parallel builds	13
8.4	Git best practices	13
9	CI/CD and Automation	14
9.1	GitHub Actions: Build, Test, Sanitize	14
9.2	Automated packaging with CPack	14

10 Debugging & Profiling	15
10.1 GDB essentials and advanced usage	15
10.2 Core dump debugging	15
10.3 Valgrind and Sanitizers	15
10.4 perf: CPU profiling	15
11 System Call and Library Tracing	17
11.1 strace and ltrace	17
11.2 ldd, ltrace, and dependency investigation	17
12 System Binary Inspection	18
12.1 ELF, symbols, and objdump	18
12.2 Static vs Dynamic linking	18
12.3 Strip and size	18
13 Containers, Sandboxing, and Reproducible Builds	19
13.1 Developer containers	19
13.2 Reproducible builds and deterministic timestamps	19
13.3 Podman vs Docker	19
14 Filesystems, Storage and Snapshots	20
14.1 Common filesystems	20
14.2 LVM basics	20
14.3 Snapshots	20
15 Packaging and Distribution	21
15.1 Creating .deb packages with CPack	21
15.2 RPM basics	21
15.3 Distribution tips	21
16 Embedded Linux and Cross-Compilation	22
16.1 Toolchain and sysroot	22
16.2 Remote debugging	22
17 Real-time Linux and Kernel Hints	23
17.1 PREEMPT_RT basics	23
17.2 Kernel build quickstart	23
17.3 iptables / nftables	23
17.4 Secure access and SSH	23
18 Monitoring, Logging and Observability	25
18.1 journalctl and logs	25
18.2 Metrics stack basics	25
19 Security Practices	26
19.1 Hardening tips	26
19.2 Secrets management	26
20 Interview Prep and Common Questions	27
20.1 Key topics to master	27
20.2 Sample questions	27

21 Mini Projects and Exercises	28
21.1 Project 1 — C++ Library with CI	28
21.2 Project 2 — Log Analyzer for Compiler Dumps	28
21.3 Project 3 — Telemetry Daemon	28
A Appendix A: Quick Reference Cheat Sheet	29
B Appendix B: 90-Day Mastery Plan (Detailed)	30
B.1 Phase 1 — Foundation (Days 1–30)	30
B.2 Phase 2 — Development (Days 31–60)	30
B.3 Phase 3 — Mastery (Days 61–90)	30
C Appendix C: Resources and Further Reading	31

Draft

1. Introduction

1.1. Why Linux Matters for C++ Developers

Linux is the standard development and deployment environment for many C++ projects, especially in systems programming, networking, embedded systems, and high-performance computing. Being comfortable with Linux significantly improves debugging capabilities, automation workflows, and overall productivity.

 **Best Practice:**

Master Linux to become a better C++ developer. Most production C++ code runs on Linux servers, and understanding the environment helps in debugging complex issues.

1.2. How to Use This Roadmap

This document is structured progressively from beginner to expert level. Each section contains practical commands, real-world examples, and industry best practices. Use the appendix as a quick reference cheat-sheet and add your personal notes where needed. Read, practice, and modify the examples — the real learning happens in the terminal.

1.3. Real-world Use Cases

- **Production Debugging:** Diagnosing crashes using `gdb` and core dumps
- **Performance Optimization:** Profiling code with `perf` and `valgrind`
- **Build Systems:** Managing complex projects with `CMake`
- **Interview Preparation:** Common Linux questions in C++ interviews
- **DevOps Integration:** CI/CD pipelines and containerization



Best Practice: Maintain a daily practice log of commands. Add frequently used commands to your shell configuration files as aliases for increased productivity.

2. Linux Basics (Beginner Level)

2.1. Filesystem Hierarchy Standard

Understanding the Linux filesystem structure is crucial:

```

1 / # Root directory
2 /bin # Essential user binaries
3 /etc # Configuration files
4 /home # User home directories
5 /usr # User programs and data
6 /var # Variable data (logs, etc.)
7 /tmp # Temporary files
8 /opt # Optional software packages
9 /dev # Device files
10 /proc # Kernel and process info (virtual FS)

```

2.2. File Operations Mastery

```

1 # Create nested directories
2 mkdir -p project/{src,include,build,test}
3
4 # Copy with progress and preserve attributes
5 rsync -av --progress source/ destination/
6
7 # Find and delete specific files
8 find . -name "*.*" -type f -delete
9
10 # Create symbolic links
11 ln -s /path/to/original link_name
12
13 # Archive and compress
14 tar -czvf project.tar.gz project/
15 tar -xzf project.tar.gz -C /tmp

```

 **Best Practice:**

Use `rsync` instead of `cp` for large directories as it's more reliable and shows progress. Use `-dry-run` first to verify operations.

2.3. Advanced Text Editing

```

1 # View multiple files
2 less file1.txt file2.txt
3
4 # Follow growing file (like logs)
5 tail -f /var/log/syslog
6
7 # View with line numbers
8 cat -n source.cpp
9
10 # Real-time log monitoring
11 tail -f application.log | ccze # if ccze is installed for coloring
12 tail -f application.log | grep -i error

```

2.4. Permission Deep Dive

```
1 # Numeric permissions explained
2 chmod 755 script.sh # rwxr-xr-x
3 chmod 644 file.txt # rw-r--r--
4 chmod 600 secret.key # rw-----
5
6 # Change ownership recursively
7 chown -R user:group project/
8
9 # Set sticky bit on directories
10 chmod +t shared_folder/
11
12 # Setuid and setgid (use with caution)
13 chmod u+s /usr/bin/somebinary
14 chmod g+s somedir/
```

 **Exercise:** Create a project structure with proper permissions and practice file operations for 30 minutes daily. Use `getfacl` / `setfacl` to experiment with ACLs.

Draft

3. Intermediate Linux Commands

3.1. Advanced File Searching

```

1 # Find C++ files modified in last 7 days
2 find . -name "*.cpp" -mtime -7
3
4 # Find files larger than 100MB
5 find / -type f -size +100M 2>/dev/null
6
7 # Find and execute commands on results
8 find . -name "*tmp" -exec rm {} \;
9
10 # Search in source code excluding build directories
11 grep -R --exclude-dir=build "TODO" .
12
13 # Use ripgrep (rg) faster, respects .gitignore
14 rg "TODO" src/ --hidden

```

3.2. System Performance Monitoring

```

1 # Real-time system monitoring
2 htop
3 glances # Comprehensive system monitor
4
5 # I/O statistics
6 iostat -x 1
7
8 # Memory usage details
9 cat /proc/meminfo
10 free -h
11
12 # Process tree view
13 pstree -p
14
15 # System resource summary
16 vmstat 1 5

```

3.3. Network Diagnostics

```

1 # Detailed network interface info
2 ip addr show
3 ip route show
4
5 # Continuous ping with timestamp
6 ping -D 8.8.8.8
7
8 # Check open ports and listening processes
9 ss -tulpn
10 netstat -tulpn
11
12 # Download with progress
13 wget --progress=bar http://example.com/file
14 curl -O http://example.com/file
15

```

```
16 # HTTP quick checks
17 curl -I https://example.com
18 curl -sS https://example.com | head -n 20
19
20 # DNS troubleshooting
21 dig example.com +short
22 nslookup example.com
```

3.4. Package Management Deep Dive

```
1 # Ubuntu/Debian advanced usage
2 sudo apt update && sudo apt upgrade
3 sudo apt install build-essential cmake git gdb valgrind
4
5 sudo apt search "c++ library"
6 sudo apt show package-name
7 sudo apt remove --purge package-name
8
9 # Snap packages
10 sudo snap install clion --classic
11
12 # Flatpak applications
13 flatpak install flathub com.jetbrains.CLion
```

 **Best Practice:** Regularly update your system and development tools. Use version pinning for critical production dependencies. Use containers to isolate toolchains when necessary.

4. Advanced System Administration

4.1. Systemd Service Management

```

1 # Create user service
2 mkdir -p ~/.config/systemd/user/
3 # Edit ~/.config/systemd/user/myservice.service with content like:
4 # [Unit]
5 # Description=My user service
6 # [Service]
7 # ExecStart=/home/user/bin/run_task.sh
8 # [Install]
9 # WantedBy=default.target
10
11 # Reload and enable
12 systemctl --user daemon-reload
13 systemctl --user enable myservice
14 systemctl --user start myservice
15
16 # View service logs
17 journalctl --user-unit=myservice -f

```

4.2. Kernel and System Optimization

```

1 # View kernel parameters
2 sysctl -a | grep tcp
3
4 # Temporary modification
5 sudo sysctl -w net.ipv4.tcp_fin_timeout=30
6
7 # Permanent changes in /etc/sysctl.conf
8 echo "net.ipv4.tcp_fin_timeout=30" | sudo tee -a /etc/sysctl.conf
9 sudo sysctl -p
10
11 # CPU governor
12 cat /sys/devices/system/cpu/cpu0/cpufreq/scaling_governor
13 sudo cpupower frequency-set -g performance

```

4.3. Advanced Security Features

```

1 # SELinux context
2 ls -Z /usr/bin/ssh
3 ps -Z
4
5 # AppArmor status
6 sudo aa-status
7
8 # Firewall management (UFW)
9 sudo ufw enable
10 sudo ufw allow 22/tcp
11 sudo ufw allow 8080/tcp
12 sudo ufw status verbose

```

4.4. Containerization with Podman/Docker

```
1 # Development container
2 docker run -it --rm -v $(pwd):/workspace \
3   -w /workspace gcc:latest bash
4
5 # Multi-stage build for C++
6 # Dockerfile example:
7 FROM gcc:latest as builder
8 WORKDIR /app
9 COPY . .
10 RUN g++ -O2 -std=c++17 -o myapp main.cpp
11
12 FROM ubuntu:latest
13 COPY --from=builder /app/myapp /usr/local/bin/myapp
14 CMD ["/usr/local/bin/myapp"]
```

 **Exercise:** Create a systemd unit that runs a small HTTP server and logs to journalctl. Test enabling, starting, stopping, and viewing logs.

Draft

5. Shell Scripting Excellence

5.1. Advanced Bash Features

```

1 #!/usr/bin/env bash
2 set -euo pipefail # Exit on error, undefined vars, pipe failures
3 IFS=$'\n\t' # Better word splitting
4
5 # Color output
6 RED='\033[0;31m'
7 GREEN='\033[0;32m'
8 NC='\033[0m' # No Color
9
10 log_success() { echo -e "${GREEN}[SUCCESS]${NC} $*"; }
11 log_error() { echo -e "${RED}[ERROR]${NC} $*" >&2; }
12
13 # Function with parameters
14 build_project() {
15     local build_type=${1:-Release}
16     log_success "Building in $build_type mode"
17     cmake -B build -DCMAKE_BUILD_TYPE=$build_type
18     cmake --build build -j$(nproc)
19 }
```

5.2. Error Handling and Debugging

```

1 # Advanced error handling
2 handle_error() {
3     local exit_code=$?
4     log_error "Command failed with exit code $exit_code"
5     # Add diagnostics: print last 20 lines of relevant log
6     journalctl -n 20 --no-pager
7     exit $exit_code
8 }
9
10 trap handle_error ERR
11
12 # Debug mode
13 if [[ "${DEBUG:-false}" == "true" ]]; then
14     set -x
15 fi
```

5.3. Configuration Management

```

1 # Sample .bashrc additions for C++ developers
2 export CC=/usr/bin/gcc
3 export CXX=/usr/bin/g++
4
5 # Development tools
6 export CMAKE_GENERATOR="Ninja"
7 export CTEST_OUTPUT_ON_FAILURE=1
8
9 # Python virtual environment
10 export WORKON_HOME=$HOME/.virtualenvs
11 export VIRTUALENVWRAPPER_PYTHON=/usr/bin/python3
```

```
12
13 # Custom functions
14 ccache_enable() { export PATH="/usr/lib/ccache:$PATH"; }
15 docker_clean() {
16     docker system prune -af
17     docker volume prune -f
18 }
```

 **Best Practice:** Place frequently used functions and aliases in a separate `/.bashrc.d/dev.sh` and source them from `/.bashrc` to keep files modular.

Draft

6. Text Processing and Data Manipulation

6.1. grep, rg, awk, sed

```

1 # grep examples
2 grep -RIn "TODO" .
3 grep -n --line-number "pattern" file
4
5 # ripgrep (rg) - fast, respects .gitignore
6 rg "TODO" --hidden
7
8 # awk examples
9 awk '{print $1, $3}' data.txt
10 awk -F: '$3 >= 1000 {print $1}' /etc/passwd
11
12 # sed examples
13 sed -n '1,100p' file.txt
14 sed -i.bak 's/old/new/g' file.txt

```

6.2. cut, sort, uniq, tr, wc

```

1 # cut
2 cut -d',' -f2 data.csv
3
4 # sort & uniq
5 sort names.txt | uniq -c | sort -nr
6
7 # tr and wc
8 tr -d '\r' < file.txt > unix_file.txt
9 wc -l *.cpp

```

6.3. Practical assembly / log parsing

```

1 # Example: find lines with register writes in assembly and list unique registers
2 rg "r[0-9]+[" asm.log | sed -E 's/.*\((r[0-9]+)\).*/\1/' | sort | uniq -c | sort -nr
3
4 # More advanced: extract destination registers from patterns like "add.f r40.x, r1.z;"
5 rg -o "r[0-9]+\.xyzw" asm.log | sed 's/\..$//' | sort | uniq -c

```

 **Project:** Write a small bash + awk pipeline to read a compiler IR dump and summarize instruction counts per opcode, output CSV for visualization.

7. Developer Productivity Tools

7.1. tmux Essentials

```
1 tmux new -s dev # start session named dev
2 tmux ls # list sessions
3 tmux attach -t dev # attach to session
4 tmux kill-session -t dev
5
6 # Split panes: prefix (Ctrl-b) then %
7 # Create window: prefix c
8 # Switch windows: prefix n / p or prefix <number>
```

7.2. fzf, ripgrep, bat

```
1 # fzf interactive file finder
2 rg --files | fzf --preview 'bat --style=numbers --color=always {}',
3
4 # Open file under cursor in vim
5 vim $(rg --files | fzf)
6
7 # bat for nice cat with syntax highlight
8 bat src/main.cpp
```

7.3. Editor setups: clangd, LSP, ccls

```
1 # clangd uses compile_commands.json generated by CMake:
2 cmake -S . -B build -DCMAKE_EXPORT_COMPILE_COMMANDS=ON
3 # In editor, point clangd to build/compile_commands.json
```

Best Practice:

Use ccache + Ninja to speed development builds and combine with clangd for fast incremental indexing.

8. Build Systems & Source Control

8.1. CMake advanced patterns

```

1 # Modern CMake example
2 cmake_minimum_required(VERSION 3.20)
3 project(MyProject LANGUAGES CXX)
4
5 set(CMAKE_CXX_STANDARD 20)
6 set(CMAKE_CXX_STANDARD_REQUIRED ON)
7
8 find_package(Boost 1.70 REQUIRED COMPONENTS filesystem system)
9
10 add_executable(myapp main.cpp)
11 target_compile_options(myapp PRIVATE -Wall -Wextra -Wpedantic)
12 target_link_libraries(myapp PRIVATE Boost::filesystem Boost::system)

```

8.2. Toolchain files and cross-compiling

```

1 # toolchain-arm.cmake (snippet)
2 set(CMAKE_SYSTEM_NAME Linux)
3 set(CMAKE_SYSTEM_PROCESSOR arm)
4 set(CMAKE_C_COMPILER arm-linux-gnueabihf-gcc)
5 set(CMAKE_CXX_COMPILER arm-linux-gnueabihf-g++)

```

8.3. Makefile tips and parallel builds

```

1 # Example wrapper Makefile for CMake
2 .PHONY: all clean format
3
4 all:
5     cmake -S . -B build -G Ninja
6     cmake --build build -- -j$(nproc)
7
8 clean:
9     rm -rf build
10
11 format:
12     clang-format -i $(shell rg --files-with-matches --glob '!build' -g '*.cpp' -g
13         '*.hpp' -S)

```

8.4. Git best practices

- Commit small, logical units. One change = one commit.
- Use descriptive commit messages (imperative tense).
- Rebase local feature branches before opening PRs to keep history linear.
- Protect main branch with required status checks in CI.

9. CI/CD and Automation

9.1. GitHub Actions: Build, Test, Sanitize

```
1 name: C++ CI
2
3 on: [push, pull_request]
4
5 jobs:
6   build:
7     runs-on: ubuntu-latest
8     strategy:
9       matrix:
10      build_type: [Debug, Release]
11
12     steps:
13       - uses: actions/checkout@v3
14       - name: Install deps
15         run: sudo apt-get update && sudo apt-get install -y build-essential cmake ninja-build
16       - name: Configure
17         run: cmake -S . -B build -G Ninja -DCMAKE_BUILD_TYPE=${{ matrix.build_type }}
18       - name: Build
19         run: cmake --build build -- -j$(nproc)
20       - name: Run tests
21         run: ctest --test-dir build --output-on-failure
```

9.2. Automated packaging with CPack

```
1 # CPack can produce .deb, .rpm, TGZ packages from CMake
2 include(InstallRequiredSystemLibraries)
3 set(CPACK_GENERATOR "DEB;TGZ")
4 set(CPACK_PACKAGE_NAME "myapp")
5 set(CPACK_PACKAGE_VERSION "1.0.0")
6 include(CPack)
```

 **Project:** Add a GitHub Action that runs clang-tidy and sanitizers on PRs. Fail the build if severity exceeds policy.

10. Debugging & Profiling

10.1. GDB essentials and advanced usage

```

1 # compile with debug info
2 g++ -g -O0 main.cpp -o main
3
4 # start gdb
5 gdb ./main
6 (gdb) break main
7 (gdb) run
8 (gdb) backtrace full
9 (gdb) info locals
10 (gdb) print variable_name
11
12 # TUI mode
13 gdb -tui ./main
14 # Reverse debugging (when enabled by gdb)
15 (gdb) record
16 (gdb) reverse-step

```

10.2. Core dump debugging

```

1 # enable core dumps
2 ulimit -c unlimited
3 # run program -> on crash core.<pid>
4 gdb ./myprog core.<pid>
5 (gdb) bt full
6 (gdb) info registers

```

10.3. Valgrind and Sanitizers

```

1 # Valgrind memcheck
2 valgrind --leak-check=full ./app
3
4 # Massif for heap profiling
5 valgrind --tool=massif ./app
6 ms_print massif.out.#####
7
8 # AddressSanitizer
9 g++ -fsanitize=address -fno-omit-frame-pointer -g main.cpp -o main
10
11 # UndefinedBehaviorSanitizer
12 g++ -fsanitize=undefined -g main.cpp -o main

```

10.4. perf: CPU profiling

```

1 # lightweight stat
2 perf stat ./app
3
4 # record call graph
5 perf record -g ./app
6 perf report -g graph

```

```
7  
8 # annotate hot functions  
9 perf annotate --stdio
```

💡 Tip:

Use ‘perf‘ for production-like sampling profiling and ‘valgrind‘/sanitizers for debug builds. Sampling profiler (perf) avoids huge overhead of instrumentation.

Draft

11. System Call and Library Tracing

11.1. strace and ltrace

```
1 # trace system calls
2 strace -f -o strace.out ./app
3
4 # summary of syscall counts and times
5 strace -c ./app
6
7 # library calls (function-level calls)
8 ltrace ./app
```

11.2. ldd, ltrace, and dependency investigation

```
1 ldd ./myprog # list dynamic dependencies
2 readelf -d myprog | less # dynamic section info
3 objdump -p myprog | grep RPATH -A 5
```

 **Exercise:** Reproduce a missing symbol error by modifying `LD_LIBRARY_PATH` and observe the `gdb/lddiagnostics`. Document the steps and fix.

12. System Binary Inspection

12.1. ELF, symbols, and objdump

```
1 readelf -h myprog # ELF header
2 readelf -s myprog | less # symbol table
3 nm -C myprog # demangled symbols
4 objdump -d myprog | less # disassembly
```

12.2. Static vs Dynamic linking

- Static linking produces a larger binary but is self-contained.
- Dynamic linking uses shared objects (.so) and reduces binary size, but requires the correct runtime environment.

12.3. Strip and size

```
1 # Strip debug/symbols to reduce size
2 strip --strip-all myprog
3 # View binary size
4 size myprog
5 file myprog
```

Draw

13. Containers, Sandboxing, and Reproducible Builds

13.1. Developer containers

```
1 # build a dev image
2 docker build -t myapp-dev -f Dockerfile.dev .
3
4 # run interactive
5 docker run --rm -it -v "$(pwd):/workspace" -w /workspace myapp-dev bash
```

13.2. Reproducible builds and deterministic timestamps

- Use fixed base images and pinned versions.
- Avoid embedding timestamps in artifacts, or set $SOURCE_DATE_EPOCH$.
- Validate by comparing checksums between builds.

13.3. Podman vs Docker

- Podman is daemonless and rootless-capable — suitable for CI on some infra.
- Most Docker files work with Podman after minor tweaks.

Draft

14. Filesystems, Storage and Snapshots

14.1. Common filesystems

- ext4 — stable default for many distros.
- xfs — high-performance, good for large files.
- btrfs — snapshots and subvolumes.
- zfs — robust data integrity (license/caveats on Linux).

14.2. LVM basics

```
1 # create physical volume, volume group, logical volume
2 pvcreate /dev/sdb1
3 vgcreate vgdata /dev/sdb1
4 lvcreate -n lvdata -L 100G vgdata
5 mkfs.xfs /dev/vgdata/lvdata
6 mount /dev/vgdata/lvdata /data
```

14.3. Snapshots

```
1 # LVM snapshot
2 lvcreate --size 1G --snapshot --name snap1 /dev/vgdata/lvdata
3 # Btrfs snapshot
4 btrfs subvolume snapshot /mnt/data /mnt/data_snapshot
```

 **Best Practice:** Test snapshot restore procedures and ensure backups are automated. Snapshots are not backups — combine with external backups.

15. Packaging and Distribution

15.1. Creating .deb packages with CPack

```
1 # In CMakeLists.txt
2 include(InstallRequiredSystemLibraries)
3 set(CPACK_GENERATOR "DEB")
4 set(CPACK_PACKAGE_NAME "myapp")
5 set(CPACK_PACKAGE_VERSION "1.0.0")
6 include(CPack)
7
8 # After build
9 cpack -G DEB -C Release
```

15.2. RPM basics

```
1 # build with rpmbuild
2 # ensure proper spec file and buildroot
3 rpmbuild -ba mypackage.spec
```

15.3. Distribution tips

- Provide clear dependency metadata.
- Use CI to produce packages for multiple distros.
- Sign packages and repositories for secure delivery.

16. Embedded Linux and Cross-Compilation

16.1. Toolchain and sysroot

```
1 export CC=arm-linux-gnueabihf-gcc
2 export CXX=arm-linux-gnueabihf-g++
3 export SYSROOT=/opt/arm/sysroot
4
5 cmake -S . -B build-arm \
6   -DCMAKE_SYSTEM_NAME=Linux \
7   -DCMAKE_SYSROOT=$SYSROOT \
8   -DCMAKE_C_COMPILER=$CC \
9   -DCMAKE_CXX_COMPILER=$CXX
```

16.2. Remote debugging

```
1 # On target:
2 gdbserver :2345 ./app
3
4 # On host:
5 arm-linux-gnueabihf-gdb ./app
6 (gdb) target remote 192.168.0.10:2345
```

 **Project:** Set up a simple cross-compilation workflow and debug a "Hello world" remotely via gdbserver. Document latency and steps.

17. Real-time Linux and Kernel Hints

17.1. PREEMPT_RT basics

```

1 # check kernel info
2 uname -a
3 # check if PREEMPT_RT is enabled (if config exposed)
4 zcat /proc/config.gz | rg PREEMPT
5
6 # test latency with cyclictest
7 sudo apt-get install rt-tests
8 sudo cyclictest -p 80 -n -i 100 -l 1000

```

17.2. Kernel build quickstart

```

1 git clone https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux-stable.git
2 cd linux-stable
3 make defconfig
4 # optionally run menuconfig to tweak options
5 make -j$(nproc)
6 # install modules (careful on production systems)
7 \end{minted}
8
9 \begin{WarningBox}
10 Compiling and installing kernels on production machines is risky. Use VM or test
     hardware. Keep recovery options (rescue kernel, bootable USB).
11 \end{WarningBox}
12
13 \clearpage
14
15 \section{Networking Deep Dive}
16
17 \subsection{Socket basics (C/C++)}
18 \begin{lstlisting}
19 // Basic TCP server skeleton (pseudocode)
20 int sock = socket(AF_INET, SOCK_STREAM, 0);
21 bind(sock, ...);
22 listen(sock, 128);
23 int client = accept(sock, ...);
24 read(client, buffer, len);

```

17.3. iptables / nftables

```

1 # List current rules
2 sudo iptables -L -n -v
3 # nftables example
4 sudo nft list ruleset

```

17.4. Secure access and SSH

```

1 ssh-keygen -t ed25519 -C "dev@company"
2 ssh-copy-id user@host
3 # Agent

```

```
4 eval "$(ssh-agent -s)"  
5 ssh-add ~/.ssh/id_ed25519
```

👉 **Exercise:** Create a small client/server pair, run server in tmux, use ss to verify listener, and simulate network failure with iptables to test reconnection logic.

Draft

18. Monitoring, Logging and Observability

18.1. journalctl and logs

```
1 # view unit logs
2 journalctl -u myservice -f
3
4 # logs since a time
5 journalctl --since "2025-01-01" --no-pager
6
7 # use rate-limit
8 journalctl --disk-usage
```

18.2. Metrics stack basics

- node_{exporter}(*Prometheus*) for host metrics.
- cAdvisor for container metrics.
- Prometheus + Grafana for alerting and dashboards.

 **Project:** Deploy a minimal Prometheus + Grafana stack using Docker Compose and create a dashboard for CPU, memory, and process counts.

Draft

19. Security Practices

19.1. Hardening tips

- Keep packages updated and apply security patches.
- Use minimal images for containers (distroless when possible).
- Run services with least privilege and dedicated users.
- Use SELinux/AppArmor and learn to debug denials.
- Use fail2ban to block brute-force attempts.

19.2. Secrets management

- Never store secrets in VCS.
- Use environment variables from secret stores or mounted volumes.
- For CI, use encrypted secrets supported by platform (GitHub Actions secrets).

Draft

20. Interview Prep and Common Questions

20.1. Key topics to master

- Process lifecycle, signals, wait/waitpid.
- Memory layout (stack, heap, BSS, data, text).
- ELF structure and dynamic linking.
- Debugging workflows and core dump analysis.
- Performance analysis: hotspots, cache behavior.
- Build systems and cross-compilation basics.

20.2. Sample questions

1. What does `set -euo pipefail` do?

Answer: `-e` exits on non-zero status, `-u` treats unset vars as error, `-o pipefail` makes whole pipeline fail if any element fails.

2. Explain `fork()` vs `exec()`.

Answer: `fork()` duplicates current process; `exec()` replaces the process image with a new program (often called in child after fork).

3. How do you debug a segmentation fault in production?

Answer: Reproduce with debug build if possible, enable core dumps, use gdb on core, analyze backtrace, inspect memory, check for UB with sanitizers.

21. Mini Projects and Exercises

21.1. Project 1 — C++ Library with CI

- Build a small library (header + implementation).
- Add unit tests with GoogleTest.
- Add CMake, CPack packaging, and GitHub Actions CI for build + tests.
- Add clang-tidy checks and sanitizer runs on PRs.

21.2. Project 2 — Log Analyzer for Compiler Dumps

- Parse compiler/assembly dumps.
- Identify registers written but not read (data flow hint).
- Output CSV with file,line,register,status.
- Visualize using simple bar charts (e.g., Python matplotlib).

21.3. Project 3 — Telemetry Daemon

- Small daemon that collects CPU, memory, disk and exposes HTTP metrics for Prometheus.
- Run as a systemd service and add health checks.
- Containerize for deployment.

Draft

A. Appendix A: Quick Reference Cheat Sheet

Command	Purpose
ls -la	List files with details and hidden files
find . -name "*.cpp"	Find C++ sources
rg "TODO"	Fast code search (ripgrep)
tmux new -s dev	Create a tmux session
gdb -q ./app	Debugging session
perf record -g ./app	Profile program
valgrind -leak-check=full ./app	Memory debugging
docker build -t myapp .	Build docker image
cmake -S . -B build	Configure CMake build
ninja -C build	Build with ninja

Draft

B. Appendix B: 90-Day Mastery Plan (Detailed)

B.1. Phase 1 — Foundation (Days 1–30)

- Week 1: Basic commands, navigation, editors.
- Week 2: Shell scripting basics, text processing.
- Week 3: Build systems (CMake), simple C++ builds.
- Week 4: Git workflows, code search tools (rg, fzf).

B.2. Phase 2 — Development (Days 31–60)

- Week 5-6: Debugging (gdb), sanitizers, valgrind.
- Week 7-8: Containers, CI, packaging, performance profiling.

B.3. Phase 3 — Mastery (Days 61–90)

- Kernel internals, system internals, real-time concepts.
- Contribute to an open-source C++ project (LLVM recommended).
- Interview prep and mock interviews.

Draft

C. Appendix C: Resources and Further Reading

- **Books:** *The Linux Programming Interface* (Michael Kerrisk), *Linux Kernel Development* (Robert Love)
- **Web:** man7.org, llvm.org, kernel.org
- **Tools:** perf, valgrind, gdb, clang-tidy, ccache, rg, fzf, tmux
- **Courses:** Explore CompilerSutra tutorials at compilersutra.com

Keep Learning, Keep Building!

CompilerSutra